

MorpheusX ASM ↔ Rust ABI Contract v1.0

Status: FROZEN

Date: 2026-01-09

Authority: Derived from NETWORK_STACK_REDESIGN.md (v1.0 Frozen) and NETWORK_STACK_AUDIT.md (v1.1)

Scope: Post-ExitBootServices bare-metal networking runtime

Target: x86_64-unknown-uefi

Document Structure

This contract is divided into mandatory chunks. Each chunk is self-contained and must be read in order. No implicit assumptions carry between chunks unless explicitly stated.

Chunk 1 — Contract Scope & Global Invariants

1.1 Purpose of This Contract

This document defines the **frozen v1 ABI** between assembly (ASM) code and Rust code for the MorpheusX post-ExitBootServices networking runtime.

This contract specifies:

- What ASM exposes to Rust
- What Rust may assume about ASM behavior
- What ASM may assume about Rust behavior
- What constitutes undefined behavior (UB)
- Memory ownership and lifetime rules
- Calling conventions and register usage

This contract does NOT specify:

- Internal ASM implementation details
- Internal Rust implementation details
- Anything outside the ASM ↔ Rust boundary

Binding Authority:

- NETWORK_STACK_REDESIGN.md §Execution Model (Locked) — execution constraints
- NETWORK_STACK_AUDIT.md Part 7 — authoritative corrections
- Where conflicts exist, AUDIT.md Part 7 supersedes REDESIGN.md

1.2 System Context (Non-Negotiable)

The following system context is immutable for v1. Any code violating these constraints exhibits undefined behavior.

Property	Value	Source
Architecture	x86_64	REDESIGN §2.1
Execution mode	Long mode (64-bit)	AUDIT §7.1.1
ABI	Microsoft x64 (UEFI-compatible)	AUDIT §7.1.6
Runtime phase	Post-ExitBootServices	REDESIGN §5.1, AUDIT §7.2.9
Interrupt state	Disabled (IF=0)	REDESIGN §2.1
Core count	1 (BSP only)	REDESIGN §Execution Model
Thread count	1	REDESIGN §Execution Model
Scheduling	None (cooperative progress)	REDESIGN §Execution Model
Execution pattern	Poll-driven	REDESIGN §Execution Model
Preemption	None	REDESIGN §Execution Model
Floating point	Disabled	REDESIGN §2.1
Async runtime	None	REDESIGN §Explicit Non-Goals
UEFI Boot Services	UNAVAILABLE	REDESIGN §5.5

1.3 Trust Boundaries

1.3.1 ASM Owns

Domain	Rationale	Source
All hardware interaction	Direct MMIO, port I/O	REDESIGN §3.1
DMA buffer submission	Descriptor ring manipulation	REDESIGN §4.5
NIC device state	Register reads/writes	REDESIGN §4.3
Memory barriers	sfence/lfence/mfence placement	AUDIT §5.4
TSC reads	rdtsc instruction	REDESIGN §7.4
Virtqueue ring updates	Available/used ring manipulation	REDESIGN §4.5
Device notification	MMIO notify writes	REDESIGN §4.5

1.3.2 Rust Owns

Domain	Rationale	Source
Protocol logic	smoltcp integration	REDESIGN §6.1
State machines	DHCP, TCP, HTTP states	REDESIGN §8.1
Application orchestration	Main loop control	REDESIGN §3.2

Domain	Rationale	Source
Buffer content preparation	Packet payload	REDESIGN §6.3
Timeout calculations	TSC delta interpretation	REDESIGN §7.4
Error handling policy	Retry vs. fatal decisions	REDESIGN §9.4

1.3.3 Boundary Crossings

The ASM ↔ Rust boundary is crossed ONLY through the functions defined in Chunk 3 (Canonical ASM Interface Table). Any other mechanism for crossing this boundary is undefined behavior.

Permitted crossing mechanisms:

- `extern "win64" fn` declarations in Rust calling ASM symbols
- ASM calling Rust functions via the `call` instruction with MS x64 ABI

Forbidden crossing mechanisms:

- Inline assembly modifying ASM-owned state
- Rust directly accessing MMIO regions
- Rust directly manipulating virtqueue rings
- Any form of shared mutable state without explicit contract

1.4 Global Invariants

These invariants are derived from REDESIGN.md and AUDIT.md and must hold at all times. Violation constitutes undefined behavior.

1.4.1 Execution Invariants

ID	Invariant	Source
EXEC-INV-1	All function calls return in bounded time	REDESIGN §2.2
EXEC-INV-2	No function may yield, await, or suspend	REDESIGN §Execution Model
EXEC-INV-3	No function may loop waiting for external state	REDESIGN §1.5 BLOCK-1
EXEC-INV-4	Interrupts remain disabled (IF=0)	REDESIGN §2.1
EXEC-INV-5	Single core execution (no SMP)	REDESIGN §Execution Model

1.4.2 Memory Invariants

ID	Invariant	Source
MEM-INV-1	DMA region is identity-mapped (<code>phys == virt</code>)	AUDIT §7.1.1
MEM-INV-2	DMA region is UC or WC (not WB)	AUDIT §7.2.2
MEM-INV-3	<code>bus_addr</code> from <code>Map()</code> used for device, <code>cpu_ptr</code> for software	AUDIT §7.2.1

ID	Invariant	Source
MEM-INV-4	Buffers have exactly one owner at any time	REDESIGN §3.3
MEM-INV-5	DEVICE-OWNED buffers are never accessed by driver	REDESIGN §3.3
MEM-INV-6	DMA region below 4GB (32-bit compatibility)	REDESIGN §4.7

1.4.3 Timing Invariants

ID	Invariant	Source
TIME-INV-1	TSC frequency calibrated at boot, never hardcoded	REDESIGN §7.6, AUDIT §7.1.4
TIME-INV-2	Invariant TSC verified via CPUID before use	REDESIGN §5.4, AUDIT §7.1.4
TIME-INV-3	TSC monotonic on single core	AUDIT §7.1.5
TIME-INV-4	wrapping_sub() for TSC comparisons	REDESIGN §7.5

1.4.4 VirtIO Invariants

ID	Invariant	Source
VIO-INV-1	Header size is 12 bytes (VERSION_1 devices)	AUDIT §7.2.7
VIO-INV-2	Feature negotiation complete before queue setup	AUDIT §7.2.4
VIO-INV-3	Queue size read from device, not hardcoded	AUDIT §7.2.5
VIO-INV-4	Device initialization only after ExitBootServices	AUDIT §7.2.9
VIO-INV-5	Descriptor table entries: 16 bytes each	VirtIO Spec 1.2 §2.6
VIO-INV-6	avail.idx increments monotonically (mod 2^16)	VirtIO Spec 1.2 §2.7

1.4.5 Barrier Invariants

ID	Invariant	Source
BAR-INV-1	sfence before avail.idx write	AUDIT §5.4.1
BAR-INV-2	mfence before device notification	AUDIT §5.4.1
BAR-INV-3	lfence after reading used.idx (or volatile)	AUDIT §5.4.2
BAR-INV-4	cli/sti are NOT memory barriers	AUDIT §7.2.11

1.5 Explicit Non-Goals (Contract Scope Exclusions)

The following are explicitly excluded from this contract. Any assumption that they are supported is incorrect.

Non-Goal	Rationale	Source
Interrupt support	Polling model	REDESIGN §Explicit Non-Goals

Non-Goal	Rationale	Source
Multi-core / SMP	Single-core design	REDESIGN §Explicit Non-Goals
Async runtime	No futures/executor	REDESIGN §Explicit Non-Goals
TLS / HTTPS	No crypto	REDESIGN §Explicit Non-Goals
IPv6	IPv4 only	REDESIGN §Explicit Non-Goals
Jumbo frames	1500 MTU only	REDESIGN §Explicit Non-Goals
Hardware offloads	Software checksums	REDESIGN §Explicit Non-Goals
UEFI Boot Services	Post-EBS only	REDESIGN §5.5
Runtime memory allocation	Pre-allocated pools only	REDESIGN §5.3
Context switching	Single execution context	REDESIGN §Execution Model

1.6 Contract Versioning

Field	Value
Contract Version	1.0
Compatible REDESIGN.md	v1.0 (Frozen)
Compatible AUDIT.md	v1.1
Breaking Change Policy	New major version required

Version Compatibility Rule: Code compiled against this contract version MUST NOT be linked with code compiled against a different major version without explicit compatibility verification.

1.7 Terminology

Term	Definition
ASM	Assembly code (x86_64 NASM/GAS syntax)
Rust	Rust code compiled for x86_64-unknown-uefi
UB	Undefined Behavior
EBS	ExitBootServices
DMA	Direct Memory Access
TSC	Time Stamp Counter
UC	Uncached memory type
WC	Write-Combining memory type
Virtqueue	VirtIO ring buffer structure
MMIO	Memory-Mapped I/O

Term	Definition
BSP	Bootstrap Processor (core 0)

Chunk 2 — Calling Convention & ABI Rules

2.1 Calling Convention: Microsoft x64 ABI

All ASM \leftrightarrow Rust boundary crossings use the **Microsoft x64 calling convention** (also known as "Win64 ABI"). This is mandated by UEFI and must be used consistently.

Source: AUDIT §7.1.6 — "On x86-64, UEFI uses the Microsoft 'Win64' calling convention"

2.1.1 Parameter Passing (Integer/Pointer)

Parameter	Register	Notes
1st	RCX	First integer/pointer argument
2nd	RDX	Second integer/pointer argument
3rd	R8	Third integer/pointer argument
4th	R9	Fourth integer/pointer argument
5th+	Stack	Right-to-left, 8-byte aligned

Stack layout for 5th+ parameters:

```
[RSP + 0x28] 5th parameter
[RSP + 0x30] 6th parameter
...
```

2.1.2 Return Values

Type	Location	Notes
Integer \leq 64 bits	RAX	Includes pointers, enums
Integer 128 bits	RAX:RDX	Low 64 in RAX, high 64 in RDX
Struct \leq 8 bytes	RAX	Returned by value
Struct $>$ 8 bytes	Hidden pointer	Caller provides buffer in RCX

2.1.3 Shadow Space Requirement

The caller MUST allocate 32 bytes of shadow space on the stack before any **call** instruction, even if fewer than 4 parameters are passed.

```
; CORRECT: Allocating shadow space before call
sub rsp, 32          ; Shadow space (minimum)
mov rcx, param1
mov rdx, param2
call target_function
add rsp, 32          ; Restore stack
```

Violation: Calling without shadow space is undefined behavior.

2.1.4 Stack Alignment

The stack MUST be 16-byte aligned **before** the `call` instruction. The `call` itself pushes 8 bytes (return address), so at function entry $RSP \bmod 16 == 8$.

Before call:	$RSP \bmod 16 == 0$	(aligned)
After call:	$RSP \bmod 16 == 8$	(misaligned by return address)
After prologue:	$RSP \bmod 16 == 0$	(realigned by push/sub)

Violation: Misaligned stack causes undefined behavior (potential crashes on SSE instructions, though we don't use SSE).

2.2 Register Classification

2.2.1 Volatile Registers (Caller-Saved)

These registers may be modified by any function call. The caller must save them if their values are needed after the call.

Register	Purpose in MS x64
RAX	Return value, scratch
RCX	1st parameter, scratch
RDX	2nd parameter, scratch
R8	3rd parameter, scratch
R9	4th parameter, scratch
R10	Scratch
R11	Scratch
RFLAGS	Volatile (except DF)
XMM0-XMM5	Floating point (unused in this contract)

2.2.2 Non-Volatile Registers (Callee-Saved)

These registers must be preserved across function calls. If a function modifies them, it must save and restore them.

Register	Notes
RBX	Must be preserved
RBP	Frame pointer (optional), must be preserved
RDI	Must be preserved
RSI	Must be preserved
RSP	Stack pointer, must be preserved
R12	Must be preserved
R13	Must be preserved
R14	Must be preserved
R15	Must be preserved
XMM6-XMM15	Must be preserved (unused in this contract)

2.2.3 Special Registers

Register	Constraint
DF (Direction Flag)	Must be 0 (clear) at call/return
IF (Interrupt Flag)	Must remain 0 (interrupts disabled)
RSP	16-byte aligned before call

2.3 Red Zone: NOT AVAILABLE

Critical: The Microsoft x64 ABI does **not** have a red zone.

Unlike the System V AMD64 ABI (which provides 128 bytes below RSP that are safe from signals), the MS x64 ABI provides no such guarantee.

Implication for ASM:

- Do NOT store data below RSP without first decrementing RSP
- All local variables must be in properly allocated stack space
- Interrupt handlers (if ever added) may clobber memory below RSP

```
; WRONG: Using non-existent red zone
mov [rsp - 8], rax ; UNSAFE - no red zone in MS ABI

; CORRECT: Properly allocate stack space
sub rsp, 8
mov [rsp], rax
```

2.4 Clobber Rules for ASM Functions

2.4.1 ASM Function Prologue Requirements

Every ASM function callable from Rust MUST:

1. Preserve all non-volatile registers it uses
2. Maintain 16-byte stack alignment
3. Allocate shadow space before calling other functions

```
; Standard ASM function prologue
my_function:
    push rbx          ; Save non-volatile (if used)
    push r12          ; Save non-volatile (if used)
    sub rsp, 40        ; 32 shadow + 8 alignment
    ; ... function body ...
    add rsp, 40
    pop r12
    pop rbx
    ret
```

2.4.2 ASM Function Epilogue Requirements

Before returning, ASM functions MUST:

1. Place return value in RAX (or RAX:RDX)
2. Restore all non-volatile registers
3. Restore RSP to entry value (minus 8 for return address)
4. Clear DF if it was modified

2.4.3 Declared Clobber Sets

For Rust `asm!` blocks calling ASM, the following clobber sets apply:

```
// Standard ASM call clobber declaration
unsafe {
    asm!(
        "call {target}",
        target = sym asm_function_name,
        // Clobbers (all volatile registers)
        out("rax") result,
        out("rcx") _,
        out("rdx") _,
        out("r8") _,
        out("r9") _,
        out("r10") _,
        out("r11") _,
        // RBX, RBP, RDI, RSI, RSP, R12-R15 preserved
```

```
    );
}
```

2.5 Floating Point and SIMD: DISABLED

Per REDESIGN §2.1: Floating point is disabled. This contract forbids:

- Use of x87 FPU instructions
- Use of SSE/AVX registers (XMM, YMM, ZMM)
- Use of any floating-point types across the boundary

Rationale: No FPU state save/restore; keeps context minimal.

Violation: Any floating-point operation across the ASM \leftrightarrow Rust boundary is undefined behavior.

2.6 Rust Declaration Syntax

Rust code declaring ASM functions must use `extern "win64"`:

```
extern "win64" {
    /// Read TSC counter
    fn asm_read_tsc() -> u64;

    /// Submit buffer to TX queue
    /// Returns: 0 on success, non-zero error code on failure
    fn asm_tx_submit(
        virtqueue_state: *mut VirtqueueState, // RCX
        buffer_idx: u16,                   // RDX (zero-extended)
        length: u16,                      // R8 (zero-extended)
    ) -> i32;
}
```

Note: On `x86_64-unknown-uefi` target, `extern "C"` is equivalent to `extern "win64"`. For clarity, this contract mandates explicit `extern "win64"`.

2.7 Parameter Type Mappings

2.7.1 Integer Types

Rust Type	Size	Register Usage	Notes
<code>u8</code>	1 byte	Low 8 bits of register	Zero-extended
<code>u16</code>	2 bytes	Low 16 bits of register	Zero-extended
<code>u32</code>	4 bytes	Low 32 bits of register	Zero-extended
<code>u64</code>	8 bytes	Full register	—
<code>usize</code>	8 bytes	Full register	Same as <code>u64</code>

Rust Type	Size	Register Usage	Notes
i8	1 byte	Low 8 bits	Sign-extended
i16	2 bytes	Low 16 bits	Sign-extended
i32	4 bytes	Low 32 bits	Sign-extended
i64	8 bytes	Full register	—
isize	8 bytes	Full register	Same as i64
bool	1 byte	Low 8 bits	0 or 1 only

2.7.2 Pointer Types

Rust Type	Size	Notes
*const T	8 bytes	Raw pointer, passed as u64
*mut T	8 bytes	Raw pointer, passed as u64
&T	8 bytes	Reference, passed as pointer
&mut T	8 bytes	Reference, passed as pointer

Pointer validity: The caller guarantees pointer validity. Null pointers are undefined behavior unless explicitly documented as permitted.

2.7.3 Struct Types

Condition	Passing Method
Size ≤ 8 bytes	By value in register
Size > 8 bytes	By hidden pointer

Contract simplification: This ABI contract does not pass structs > 8 bytes across the boundary. All complex data is passed via pointers to pre-defined structures.

2.8 Function Attributes

2.8.1 Required Attributes for ASM Symbols

ASM functions must be declared with:

- Global visibility (.global or global)
- Proper alignment (function entry 16-byte aligned recommended)
- No stack protector dependencies

2.8.2 Required Attributes for Rust FFI

```
#[no_mangle] // If Rust function called from ASM
extern "win64" fn rust_callback(...) { }
```

2.9 Calling Convention Summary Table

Aspect	Requirement	Violation Consequence
Parameter registers	RCX, RDX, R8, R9	Wrong values passed
Return register	RAX (or RAX:RDX)	Wrong return value
Shadow space	32 bytes minimum	Stack corruption
Stack alignment	16-byte before call	Potential crash
Non-volatile preservation	RBX, RBP, RDI, RSI, R12-R15	Caller state corruption
Direction flag	Clear (DF=0)	String ops malfunction
Red zone	None (forbidden)	Stack corruption

Chunk 3 — Canonical ASM Interface Table

3.1 Interface Completeness Guarantee

This section defines the **complete** set of ASM-exported symbols that Rust may call. Any symbol not listed here is internal to ASM and must not be called from Rust.

Invariant: The $\text{ASM} \leftrightarrow \text{Rust}$ interface is CLOSED. No additional functions may be added without a contract revision.

3.2 Symbol Naming Convention

All ASM symbols callable from Rust follow this pattern:

```
asm_<subsystem>_<operation>
```

Subsystem	Purpose
tsc	Time Stamp Counter operations
vq	Virtqueue operations
mmio	Memory-mapped I/O
bar	Memory barrier operations
nic	NIC-specific operations

3.3 Complete ASM Interface Table

3.3.1 Time Stamp Counter Operations

asm_tsc_read

Symbol:	asm_tsc_read
Purpose:	Read the current TSC value
Inputs:	None
Outputs:	RAX = 64-bit TSC value
Clobbers:	RAX, RDX (rdtsc clobbers EDX:EAX)
Barriers:	None (not serializing)
Errors:	None (always succeeds)
Latency:	~20-40 cycles

Rust declaration:

```
extern "win64" {
    fn asm_tsc_read() -> u64;
}
```

ASM implementation requirements:

```
asm_tsc_read:
    rdtsc                      ; EDX:EAX = TSC
    shl rdx, 32                 ; RAX = full 64-bit TSC
    or  rax, rdx                ; RAX = full 64-bit TSC
    ret
```

asm_tsc_read_serialized

Symbol:	asm_tsc_read_serialized
Purpose:	Read TSC with serialization (for precise measurement)
Inputs:	None
Outputs:	RAX = 64-bit TSC value
Clobbers:	RAX, RBX, RCX, RDX (cpuid + rdtsc)
Barriers:	Full serialization via CPUID
Errors:	None
Latency:	~100-200 cycles (CPUID overhead)

Rust declaration:

```
extern "win64" {
    fn asm_tsc_read_serialized() -> u64;
}
```

ASM implementation requirements:

```
asm_tsc_read_serialized:
    push rbx                      ; CPUID clobbers RBX (non-volatile)
    xor eax, eax
    cpuid                         ; Serialize
    rdtsc
    shl rdx, 32
    or rax, rdx
    pop rbx
    ret
```

3.3.2 Virtqueue Operations

asm_vq_submit_tx

Symbol: asm_vq_submit_tx
 Purpose: Submit a buffer to the TX virtqueue with proper barriers
 Inputs:
 RCX = *mut VirtqueueState (pointer to TX queue state)
 RDX = buffer_index: u16 (descriptor index to submit)
 R8 = buffer_length: u16 (total length including 12-byte header)
 Outputs:
 RAX = 0 on success
 RAX = 1 if queue full
 Clobbers: RAX, RCX, RDX, R8, R9, R10, R11
 Barriers:
 - sfence after descriptor write
 - sfence after avail.ring write
 - mfence before notify (if notify performed)
 Errors: Returns 1 if no free descriptors
 Latency: ~50-100 cycles (barrier-dependent)

Rust declaration:

```
extern "win64" {
    fn asm_vq_submit_tx(
        vq_state: *mut VirtqueueState,
        buffer_index: u16,
        buffer_length: u16,
    ) -> u32;
}
```

Preconditions (Rust must guarantee):

- `vq_state` points to valid, initialized VirtqueueState
- `buffer_index` < `queue_size`
- Buffer at `buffer_index` is DRIVER-OWNED
- Buffer content already written (including 12-byte VirtIO header)

Postconditions (ASM guarantees):

- On success (0): Buffer is now DEVICE-OWNED
- On failure (1): Buffer ownership unchanged

asm_vq_poll_tx_complete

Symbol: `asm_vq_poll_tx_complete`
 Purpose: Poll TX used ring for completed transmissions
 Inputs:
 `RCX` = `*mut VirtqueueState` (pointer to TX queue state)
 Outputs:
 `RAX` = `buffer_index`: `u32` (low 16 bits) if completion found
 `RAX` = `0xFFFFFFFF` if no completion
 Clobbers: `RAX, RCX, RDX, R8, R9, R10, R11`
 Barriers:
 - lfence after reading `used.idx` (or volatile read)
 - lfence before reading buffer index from used ring
 Errors: None (empty queue returns sentinel)
 Latency: ~30-50 cycles

Rust declaration:

```
extern "win64" {
    fn asm_vq_poll_tx_complete(vq_state: *mut VirtqueueState) -> u32;
}
```

Postconditions:

- If return != `0xFFFFFFFF`: Buffer at returned index is now DRIVER-OWNED
- Caller must return buffer to free pool or reuse

asm_vq_submit_rx

Symbol: `asm_vq_submit_rx`
 Purpose: Submit an empty buffer to RX virtqueue for receiving
 Inputs:
 `RCX` = `*mut VirtqueueState` (pointer to RX queue state)

```

RDX = buffer_index: u16 (descriptor index)
R8  = buffer_capacity: u16 (buffer size, must be  $\geq$  12 + MTU)
Outputs:
    RAX = 0 on success
    RAX = 1 if queue full
Clobbers:   RAX, RCX, RDX, R8, R9, R10, R11
Barriers:   Same as asm_vq_submit_tx
Errors:     Returns 1 if no free descriptors

```

Rust declaration:

```

extern "win64" {
    fn asm_vq_submit_rx(
        vq_state: *mut VirtqueueState,
        buffer_index: u16,
        buffer_capacity: u16,
    ) -> u32;
}

```

Preconditions:

- Buffer at `buffer_index` is DRIVER-OWNED
- `buffer_capacity` \geq 1526 (12-byte header + 1514-byte max frame)

asm_vq_poll_rx

```

Symbol:      asm_vq_poll_rx
Purpose:    Poll RX used ring for received packets
Inputs:
    RCX = *mut VirtqueueState (pointer to RX queue state)
    RDX = *mut RxResult (output structure for result)
Outputs:
    RAX = 0 if no packet available
    RAX = 1 if packet received (result written to RDX)
Clobbers:   RAX, RCX, RDX, R8, R9, R10, R11
Barriers:
    - lfence after reading used.idx
    - lfence before reading packet data
Errors:     None

```

RxResult structure:

```

#[repr(C)]
pub struct RxResult {
    pub buffer_index: u16,    // Which buffer received the packet
    pub length: u16,          // Bytes written (including 12-byte header)
}

```

```
    pub _reserved: u32,           // Padding for alignment
}
```

Rust declaration:

```
extern "win64" {
    fn asm_vq_poll_rx(
        vq_state: *mut VirtqueueState,
        result: *mut RxResult,
    ) -> u32;
}
```

Postconditions:

- If return == 1: Buffer at `result.buffer_index` is DRIVER-OWNED
- Caller must process packet, then resubmit buffer or free

asm_vq_notify

Symbol: `asm_vq_notify`
 Purpose: Notify device that new buffers are available
 Inputs:
 RCX = `*mut VirtqueueState` (pointer to queue state)
 Outputs: None (`void`)
 Clobbers: RAX, RCX, RDX, R8, R9, R10, R11
 Barriers: mfence before MMIO write (implicit in implementation)
 Errors: None
 Latency: ~100-500 cycles (MMIO write)

Rust declaration:

```
extern "win64" {
    fn asm_vq_notify(vq_state: *mut VirtqueueState);
}
```

Notes:

- Notification may be suppressed if device indicates no notification needed
- Suppression check is internal to ASM

3.3.3 Memory Barrier Operations**asm_bar_sfence**

Symbol:	asm_bar_sfence
Purpose:	Store fence - ensures all prior stores are visible
Inputs:	None
Outputs:	None
Clobbers:	None
Barriers:	sfence

Rust declaration:

```
extern "win64" {
    fn asm_bar_sfence();
}
```

asm_bar_lfence

Symbol:	asm_bar_lfence
Purpose:	Load fence - ensures all prior loads complete
Inputs:	None
Outputs:	None
Clobbers:	None
Barriers:	lfence

Rust declaration:

```
extern "win64" {
    fn asm_bar_lfence();
}
```

asm_bar_mfence

Symbol:	asm_bar_mfence
Purpose:	Full fence - ensures all prior loads AND stores complete
Inputs:	None
Outputs:	None
Clobbers:	None
Barriers:	mfence

Rust declaration:

```
extern "win64" {
    fn asm_bar_mfence();
```

```
}
```

3.3.4 MMIO Operations

asm_mmio_read32

```
Symbol:      asm_mmio_read32
Purpose:     Read 32-bit value from MMIO address
Inputs:
    RCX = address: u64 (MMIO address, must be 4-byte aligned)
Outputs:
    RAX = value: u32 (zero-extended to 64 bits)
Clobbers:   RAX
Barriers:   Implicit (MMIO on x86 is UC/strongly ordered)
```

Rust declaration:

```
extern "win64" {
    fn asm_mmio_read32(address: u64) -> u32;
}
```

asm_mmio_write32

```
Symbol:      asm_mmio_write32
Purpose:     Write 32-bit value to MMIO address
Inputs:
    RCX = address: u64 (MMIO address, must be 4-byte aligned)
    RDX = value: u32
Outputs:     None
Clobbers:   None
Barriers:   Implicit
```

Rust declaration:

```
extern "win64" {
    fn asm_mmio_write32(address: u64, value: u32);
}
```

asm_mmio_read16 / asm_mmio_write16

Symbol: asm_mmio_read16 / asm_mmio_write16
 Purpose: 16-bit MMIO read/write
 Inputs/Outputs: Same pattern as 32-bit variants
 Alignment: 2-byte aligned

3.3.5 NIC Initialization

asm_nic_reset

Symbol: asm_nic_reset
 Purpose: Reset VirtIO NIC device
 Inputs:
 RCX = mmio_base: u64 (device MMIO base address)
 Outputs:
 RAX = 0 on success
 RAX = 1 on timeout (device did not reset)
 Clobbers: RAX, RCX, RDX, R8, R9, R10, R11
 Barriers: Internal barriers as needed
 Timeout: 100ms maximum

Rust declaration:

```
extern "win64" {
    fn asm_nic_reset(mmio_base: u64) -> u32;
}
```

asm_nic_set_status

Symbol: asm_nic_set_status
 Purpose: Write device status register
 Inputs:
 RCX = mmio_base: u64
 RDX = status: u8 (status bits to write)
 Outputs: None
 Clobbers: RAX, RCX, RDX

asm_nic_get_status

Symbol: asm_nic_get_status
 Purpose: Read device status register
 Inputs:
 RCX = mmio_base: u64
 Outputs:

```
RAX = status: u8 (current status bits)
Clobbers: RAX
```

asm_nic_read_features

```
Symbol: asm_nic_read_features
Purpose: Read device feature bits
Inputs:
    RCX = mmio_base: u64
Outputs:
    RAX = features: u64 (device-offered features)
Clobbers: RAX, RCX, RDX
```

asm_nic_write_features

```
Symbol: asm_nic_write_features
Purpose: Write driver-accepted feature bits
Inputs:
    RCX = mmio_base: u64
    RDX = features: u64 (features to accept)
Outputs: None
Clobbers: RAX, RCX, RDX
```

asm_nic_read_mac

```
Symbol: asm_nic_read_mac
Purpose: Read MAC address from device config space
Inputs:
    RCX = mmio_base: u64
    RDX = *mut [u8; 6] (output buffer for MAC)
Outputs:
    RAX = 0 on success
    RAX = 1 if MAC not available (feature not negotiated)
Clobbers: RAX, RCX, RDX, R8
```

Rust declaration:

```
extern "win64" {
    fn asm_nic_read_mac(mmio_base: u64, mac_out: *mut [u8; 6]) -> u32;
}
```

3.4 Shared Data Structures

3.4.1 VirtqueueState

```
/// Virtqueue state shared between ASM and Rust.  
/// This structure MUST NOT be modified by Rust except at initialization.  
#[repr(C, align(64))] // Cache line aligned  
pub struct VirtqueueState {  
    // === Descriptor Table ===  
    /// Physical address of descriptor table  
    pub desc_table_phys: u64,  
    /// CPU pointer to descriptor table  
    pub desc_table_ptr: *mut VirtqDesc,  
  
    // === Available Ring ===  
    /// Physical address of available ring  
    pub avail_ring_phys: u64,  
    /// CPU pointer to available ring  
    pub avail_ring_ptr: *mut VirtqAvail,  
  
    // === Used Ring ===  
    /// Physical address of used ring  
    pub used_ring_phys: u64,  
    /// CPU pointer to used ring  
    pub used_ring_ptr: *mut VirtqUsed,  
  
    // === Queue Parameters ===  
    /// Number of entries (power of 2, ≤ 32768)  
    pub queue_size: u16,  
    /// Queue index (0 = RX, 1 = TX)  
    pub queue_index: u16,  
    /// Padding  
    pub _pad0: u32,  
  
    // === Driver State (ASM-owned) ===  
    /// Next descriptor index to use  
    pub next_desc_idx: u16,  
    /// Last seen used index  
    pub last_used_idx: u16,  
    /// Number of free descriptors  
    pub free_desc_count: u16,  
    /// Padding  
    pub _pad1: u16,  
  
    // === Device Notification ===  
    /// MMIO address for queue notification  
    pub notify_addr: u64,  
    /// Notification offset (for modern VirtIO)  
    pub notify_offset: u32,  
    /// Padding  
    pub _pad2: u32,  
  
    // === Buffer Tracking ===  
    /// Pointer to buffer physical addresses array
```

```

pub buffer_phys_addrs: *const u64,
/// Pointer to buffer CPU addresses array
pub buffer_cpu_addrs: *const *mut u8,
/// Size of each buffer in bytes
pub buffer_size: u32,
/// Number of buffers
pub buffer_count: u32,
}

```

3.4.2 VirtIO Descriptor

```

/// VirtIO descriptor table entry (16 bytes)
#[repr(C)]
pub struct VirtqDesc {
    /// Physical address of buffer
    pub addr: u64,
    /// Length of buffer
    pub len: u32,
    /// Flags (NEXT, WRITE, INDIRECT)
    pub flags: u16,
    /// Next descriptor index (if NEXT flag set)
    pub next: u16,
}

```

3.4.3 VirtIO Available Ring

```

/// VirtIO available ring
#[repr(C)]
pub struct VirtqAvail {
    /// Flags (notification suppression)
    pub flags: u16,
    /// Next available index (incremented by driver)
    pub idx: u16,
    /// Ring entries (variable size, queue_size entries)
    pub ring: [u16; 0], // Flexible array member
    // Followed by: used_event: u16 (if VIRTIO_F_EVENT_IDX)
}

```

3.4.4 VirtIO Used Ring

```

/// VirtIO used ring
#[repr(C)]
pub struct VirtqUsed {
    /// Flags (notification suppression)
    pub flags: u16,
    /// Next used index (incremented by device)
}

```

```

pub idx: u16,
    /// Ring entries (variable size, queue_size entries)
pub ring: [VirtqUsedElem; 0], // Flexible array member
    // Followed by: avail_event: u16 (if VIRTIO_F_EVENT_IDX)
}

#[repr(C)]
pub struct VirtqUsedElem {
    /// Descriptor chain head index
    pub id: u32,
    /// Bytes written by device
    pub len: u32,
}

```

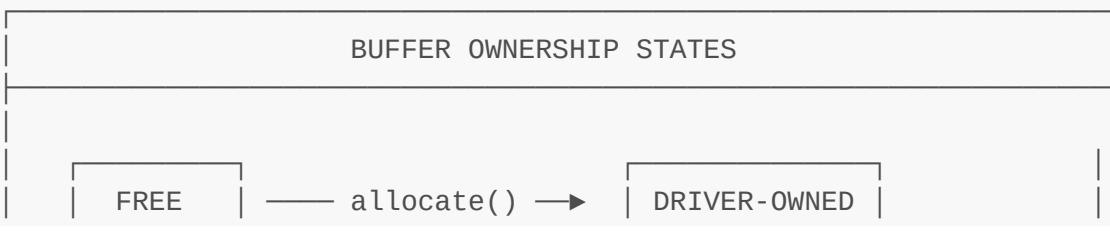
3.5 Interface Constraints Summary

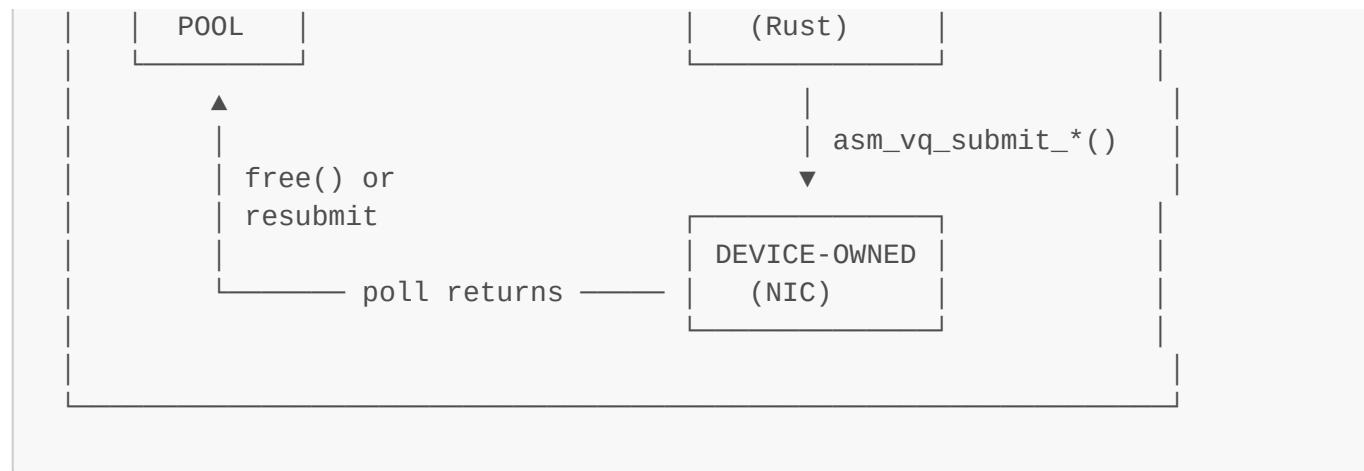
Symbol	Max Latency	May Block	Barriers
asm_tsc_read	40 cycles	No	None
asm_tsc_read_serialized	200 cycles	No	CPUID
asm_vq_submit_tx	100 cycles	No	sfence, mfence
asm_vq_poll_tx_complete	50 cycles	No	lfence
asm_vq_submit_rx	100 cycles	No	sfence, mfence
asm_vq_poll_rx	50 cycles	No	lfence
asm_vq_notify	500 cycles	No	mfence
asm_bar_*	50 cycles	No	As named
asm_mmio_*	200 cycles	No	Implicit
asm_nic_reset	100ms	Yes (bounded)	Internal

Chunk 4 — Ownership & Lifetime Semantics

4.1 Ownership Model Overview

Buffer ownership in this ABI follows an **exclusive ownership** model. At any instant, each DMA buffer is in exactly one of three states:





Source: REDESIGN §3.3 Buffer Ownership Model

4.2 State Definitions

4.2.1 FREE State

A buffer is FREE when:

- It is in the allocator's free list
- No descriptor references it
- Neither driver nor device may access its contents

Transitions:

- FREE → DRIVER-OWNED: Via buffer allocation

4.2.2 DRIVER-OWNED State

A buffer is DRIVER-OWNED when:

- It has been allocated but not submitted to a virtqueue
- OR it has been returned from the used ring but not freed

Permitted operations (Rust):

- Read buffer contents
- Write buffer contents
- Free buffer (return to FREE)
- Submit to virtqueue (transition to DEVICE-OWNED)

Forbidden operations (Rust):

- None (full access permitted)

4.2.3 DEVICE-OWNED State

A buffer is DEVICE-OWNED when:

- It has been submitted to a virtqueue via `asm_vq_submit_*`
- AND it has not yet appeared in the used ring

Permitted operations:

- None from Rust
- Device (NIC) may read (TX) or write (RX) the buffer

Forbidden operations (Rust):

- Read buffer contents
- Write buffer contents
- Free buffer
- Resubmit to any queue

Violation: Any access to a DEVICE-OWNED buffer from Rust is **undefined behavior**.

4.3 Ownership Transfer Rules

4.3.1 Rule OWN-1: Submission Transfers Ownership to Device

RULE OWN-1: Virtqueue Submission

PRECONDITION:

Buffer B is DRIVER-OWNED
Buffer content is fully initialized (TX) or ready to receive (RX)

ACTION:

Call `asm_vq_submit_tx(state, B.index, length)`
OR `asm_vq_submit_rx(state, B.index, capacity)`

POSTCONDITION (if success):

Buffer B is DEVICE-OWNED
Rust MUST NOT access B until poll returns it

POSTCONDITION (if failure - queue full):

Buffer B remains DRIVER-OWNED
Rust may retry or take other action

Source: REDESIGN §3.3 DMA-OWN-1, AUDIT §2.1

4.3.2 Rule OWN-2: Poll Completion Transfers Ownership to Driver

RULE OWN-2: Virtqueue Completion

PRECONDITION:

Buffer B is DEVICE-OWNED (submitted earlier)
Device has placed B's descriptor index in used ring

ACTION:

Call `asm_vq_poll_tx_complete(state) → returns B.index`
OR `asm_vq_poll_rx(state, result) → returns B.index in result`

POSTCONDITION:

- Buffer B is DRIVER-OWNED
- Rust may access B.contents
- For RX: B contains packet data (12-byte header + frame)
- For TX: B may be freed or reused

Source: REDESIGN §3.3 DMA-OWN-2, DMA-OWN-3

4.3.3 Rule OWN-3: Free Returns to Pool

RULE OWN-3: Buffer Deallocation**PRECONDITION:**

- Buffer B is DRIVER-OWNED
- No references to B exist

ACTION:

- Return B to free pool (Rust-side allocator)

POSTCONDITION:

- Buffer B is FREE
- B may be reallocated for different purpose

4.4 RX Buffer Lifecycle

LIFECYCLE RX-1: Receive Buffer**1. ALLOCATE**

- Rust: B = alloc_rx_buffer()
- State: B is DRIVER-OWNED

2. SUBMIT

- Rust: asm_vq_submit_rx(state, B.index, capacity)
- State: B is DEVICE-OWNED
- Rust: MUST NOT access B

3. RECEIVE (async, by device)

- Device: Writes packet to B
- Device: Marks B in used ring with length
- State: B is still DEVICE-OWNED (from Rust perspective)

4. POLL

- Rust: asm_vq_poll_rx(state, result) → returns 1
- State: B is DRIVER-OWNED
- result.buffer_index = B.index
- result.length = bytes written (including 12-byte header)

5. CONSUME

- Rust: Process packet in B[12..result.length]

B[0..12] contains VirtIO header (typically ignored)

6. RECYCLE or FREE

Option A: `asm_vq_submit_rx(state, B.index, capacity) → goto step 2`

Option B: `free_rx_buffer(B) → B is FREE`

Timing constraint: Steps 4-6 should complete within one main loop iteration to prevent RX queue starvation.

4.5 TX Buffer Lifecycle

LIFECYCLE TX-1: Transmit Buffer

1. ALLOCATE

Rust: `B = alloc_tx_buffer()`

State: `B is DRIVER-OWNED`

2. PREPARE

Rust: Zero `B[0..12]` (VirtIO header, no offloads)

Rust: Copy Ethernet frame to `B[12..12+frame_len]`

Total length = $12 + \text{frame_len}$

3. SUBMIT

Rust: `asm_vq_submit_tx(state, B.index, 12 + frame_len)`

State: `B is DEVICE-OWNED`

Rust: MUST NOT access `B`

4. TRANSMIT (async, by device)

Device: Reads packet from `B`

Device: Transmits on wire

Device: Marks `B` in used ring

State: `B is still DEVICE-OWNED (from Rust perspective)`

5. POLL COMPLETION

Rust: `asm_vq_poll_tx_complete(state) → returns B.index`

State: `B is DRIVER-OWNED`

6. FREE

Rust: `free_tx_buffer(B)`

State: `B is FREE`

Note: TX completion may be delayed. Rust must not assume completion timing.

4.6 smoltcp Token Lifetimes

smoltcp's Device trait uses tokens to manage buffer lifetimes. This section specifies how tokens interact with ASM ownership.

4.6.1 RxToken Lifetime

```
pub struct VirtioRxToken<'a> {
    driver: &'a mut VirtioDriver,
    buffer_index: u16,
    length: u16,
}
```

Lifetime rule:

- RxToken borrows `&mut VirtioDriver`
- While RxToken exists, buffer is DRIVER-OWNED
- `consume()` must be called before RxToken is dropped
- After `consume()`, buffer may be resubmitted

INVARIANT TOKEN-RX-1: RxToken Buffer Validity

WHILE RxToken exists:

- Buffer at token.buffer_index is DRIVER-OWNED
- Buffer contents are valid and readable

ON RxToken::consume():

- Closure F receives `&mut buffer[12..length]`
- After F returns, buffer may be resubmitted

ON RxToken::drop() without consume():

- Buffer must still be resubmitted or freed
- Packet data is lost (acceptable)

4.6.2 TxToken Lifetime

```
pub struct VirtioTxToken<'a> {
    driver: &'a mut VirtioDriver,
}
```

Lifetime rule:

- TxToken borrows `&mut VirtioDriver`
- `consume(len, F)` allocates buffer, calls F, then submits
- After `consume()`, buffer is DEVICE-OWNED

INVARIANT TOKEN-TX-1: TxToken Submission

ON TxToken::consume(len, F):

1. Allocate buffer B (DRIVER-OWNED)
2. Zero B[0..12] (VirtIO header)
3. Call F(&mut B[12..12+len])
4. Call `asm_vq_submit_tx(state, B.index, 12+len)`

5. B is now DEVICE-OWNED

```
ON TxToken::drop() without consume():
    No buffer allocated, no action needed
```

4.7 Descriptor Table Ownership

The descriptor table is shared between ASM and Rust with specific ownership rules.

4.7.1 Descriptor Lifecycle

DESCRIPTOR STATE MACHINE:

```
FREE → DRIVER-OWNED → IN-FLIGHT → DRIVER-OWNED → FREE
      (preparing)          (submitted)        (completed)
```

FREE:

- Descriptor in ASM free list
- ASM may allocate

DRIVER-OWNED (preparing):

- ASM has allocated descriptor
- Rust fills in addr, len, flags via VirtqueueState pointers
- Not yet in available ring

IN-FLIGHT:

- Descriptor index in available ring
- Device may read descriptor
- ASM tracks in in_flight_count

DRIVER-OWNED (completed):

- Descriptor index returned from used ring
- Rust may read completion status
- ASM returns to free list

4.7.2 Descriptor Access Rules

Field	Who May Write	When
desc.addr	Rust (via ASM init)	Before submit only
desc.len	Rust (via ASM init)	Before submit only
desc.flags	Rust (via ASM init)	Before submit only
desc.next	ASM	Always (free list management)
avail.idx	ASM	On submit
avail.ring[i]	ASM	On submit

Field	Who May Write	When
<code>used.idx</code>	Device	Async
<code>used.ring[i]</code>	Device	Async

4.8 Lifetime Invariants Summary

ID	Invariant	Violation Consequence
LIFE-1	Buffer in exactly one state at any time	Data corruption
LIFE-2	DEVICE-OWNED buffers never accessed by Rust	Data corruption, UB
LIFE-3	RxToken consume() before resubmit	Use-after-free
LIFE-4	TxToken consume() transfers ownership	Buffer leak
LIFE-5	Descriptors freed only after used ring return	Descriptor leak
LIFE-6	Buffer physical address stable for lifetime	DMA corruption

4.9 Buffer Memory Requirements

4.9.1 Buffer Sizing

Buffer Type	Minimum Size	Recommended	Contents
RX	1526 bytes	2048 bytes	12B header + 1514B max frame
TX	1526 bytes	2048 bytes	12B header + 1514B max frame

4.9.2 Buffer Alignment

Requirement	Value	Rationale
Start alignment	8 bytes	Descriptor addr field
DMA alignment	4 bytes minimum	VirtIO spec
Recommended	64 bytes	Cache line

4.9.3 Buffer Count

Queue	Minimum	Recommended	Maximum
RX	2	32	queue_size
TX	2	32	queue_size

Constraint: Buffer count \leq queue_size (negotiated with device).

Chunk 5 — Memory Ordering & DMA Visibility

5.1 Memory Model Context

This ABI operates under the x86-64 Total Store Order (TSO) memory model with the following characteristics:

Property	x86-64 Behavior	Source
Store-Store ordering	Preserved	Intel SDM §8.2.2
Load-Load ordering	Preserved	Intel SDM §8.2.2
Load-Store ordering	Preserved	Intel SDM §8.2.2
Store-Load ordering	MAY reorder	Intel SDM §8.2.2
Compiler reordering	Must be prevented	Volatile/barriers

Critical: The TSO model applies to CPU ↔ CPU communication. DMA (CPU ↔ Device) has additional requirements.

Source: AUDIT §2.3, §7.2.2

5.2 The Three Ordering Concerns

This contract distinguishes three separate ordering concerns:

5.2.1 Compiler Ordering

Problem: The Rust/LLVM compiler may reorder memory accesses.

Solution: Use `volatile` operations or inline assembly barriers.

```
// Compiler barrier (does NOT emit CPU instruction)
core::sync::atomic::compiler_fence(Ordering::SeqCst);

// Volatile read (prevents compiler reordering)
core::ptr::read_volatile(addr);

// Volatile write (prevents compiler reordering)
core::ptr::write_volatile(addr, value);
```

Contract rule: All MMIO and DMA buffer accesses must use volatile operations or be inside ASM functions.

5.2.2 CPU Store Buffer Ordering

Problem: CPU stores may sit in the store buffer, invisible to other agents.

Solution: Use `sfence` (store fence) to flush store buffer.

```
; After writing to DMA buffer:  
sfence           ; Ensure stores visible to coherency domain
```

Contract rule: ASM inserts `sfence` before any operation that makes data visible to device.

5.2.3 Cache Coherency

Problem: On x86 with standard write-back memory, CPU caches are coherent with DMA by hardware design. However, this requires:

- The memory region be in the coherency domain
- No non-coherent memory types (with some exceptions)

Solution (this contract): DMA regions mapped as Uncached (UC) or Write-Combining (WC).

Source: AUDIT §7.2.2 — "DMA region MUST be one of: a) Uncached (UC), b) Write-Combining (WC), c) Write-Back (WB) with explicit cache management"

Contract choice: UC or WC mapping. No explicit `clflush` required.

5.3 Barrier Placement Contracts

5.3.1 TX Submission Barrier Contract

CONTRACT MEM-TX-1: TX Buffer Submission Ordering

SEQUENCE (must execute in this exact order):

1. WRITE buffer data
 - Rust: Writes packet to `buffer[0..length]`
 - Memory: UC/WC, so writes go directly to memory (no cache)
 - No barrier needed between buffer writes (TSO)
2. WRITE descriptor
 - ASM: Writes `desc.addr`, `desc.len`, `desc.flags`
 - Memory: UC mapping for descriptor table
3. SFENCE
 - ASM: `sfence` instruction
 - Purpose: Ensure descriptor visible before index update
4. WRITE available ring entry
 - ASM: `avail.ring[avail.idx & mask] = desc_index`
5. SFENCE
 - ASM: `sfence` instruction
 - Purpose: Ensure ring entry visible before index update
6. WRITE available index
 - ASM: `avail.idx += 1`

7. MFENCE
 - ASM: mfence instruction
 - Purpose: Full barrier before notify (store-load ordering)
8. CONDITIONAL NOTIFY
 - ASM: Check if notification needed
 - ASM: If yes, write to notify MMIO register

RATIONALE:

Device may read descriptors immediately after seeing avail.idx change. Without barriers, device may see stale descriptor data. mfence before notify prevents notify being observed before idx update.

Source: AUDIT §5.4.1

5.3.2 RX Completion Barrier Contract

CONTRACT MEM-RX-1: RX Completion Polling Ordering

SEQUENCE (must execute in this exact order):

1. READ used.idx
 - ASM: Volatile read of used.idx
 - Note: x86 provides acquire semantics on reads to UC memory
2. COMPARE
 - ASM: if used.idx == last_used_idx, return (no completions)
3. LFENCE
 - ASM: lfence instruction
 - Purpose: Ensure idx read completes before ring entry read
 - Note: May be unnecessary on x86 TSO, but required for correctness proof
4. READ used ring entry
 - ASM: entry = used.ring[last_used_idx & mask]
 - Obtains: desc_index and bytes_written
5. LFENCE (or implicit via data dependency)
 - ASM: lfence before reading buffer contents
 - Purpose: Ensure ring read completes before buffer access
6. RETURN to caller
 - ASM: Returns desc_index and length
 - Rust: May now access buffer contents
7. UPDATE tracking
 - ASM: last_used_idx += 1 (internal bookkeeping)

RATIONALE:

Device writes buffer → writes used ring → increments used.idx
 Driver must read in reverse order with barriers to see consistent data.

Source: AUDIT §5.4.2

5.3.3 Device Notification Barrier Contract

CONTRACT MEM-NOTIFY-1: Device Notification

PRECONDITION:

All relevant available ring updates complete
 avail.idx reflects all submitted descriptors

SEQUENCE:

1. MFENCE
 - Purpose: Ensure all stores visible before MMIO write
2. CHECK notification suppression (optional optimization)
 - Read used.flags for VIRTQ_USED_F_NO_NOTIFY
 - If set, skip notification
3. MMIO WRITE to notify register
 - VirtIO modern: queue_notify register
 - This is a posted write (no completion)

RATIONALE:

MMIO writes are UC and serialize at the bus level.
 mfence ensures all prior DMA memory writes are visible
 before the device processes the notification.

5.4 Memory Type Requirements

5.4.1 DMA Region Memory Type

Region	Required Type	Alternative	Source
Descriptor tables	UC	WC	AUDIT §2.4 CACHE-2
Available rings	UC	WC	AUDIT §2.4 CACHE-2
Used rings	UC	WC	AUDIT §2.4 CACHE-2
RX buffers	UC	WC	AUDIT §2.4 CACHE-1
TX buffers	UC	WC (preferred)	AUDIT §2.4 CACHE-1

UC (Uncached):

- Reads/writes go directly to memory
- No caching, no speculation

- Slowest but simplest

WC (Write-Combining):

- Writes may be combined in WC buffer
- Reads go directly to memory
- Better performance for sequential writes (TX buffers)
- Requires `sfence` to flush WC buffers

5.4.2 MMIO Memory Type

All MMIO regions (device registers) are implicitly UC on x86.

Region	Type	Ordering
VirtIO capability registers	UC	Strongly ordered
Notify region	UC	Strongly ordered
Device config space	UC	Strongly ordered

5.5 Visibility Guarantees

5.5.1 Rust → Device Visibility

GUARANTEE VIS-1: Rust Write Visibility to Device

GIVEN:

Rust writes data to DMA buffer (DRIVER-OWNED)
Rust calls `asm_vq_submit_tx()`

THEN:

Device will see the written data

MECHANISM:

1. Buffer is UC/WC → writes bypass cache
2. ASM inserts `sfence` before `avail.idx` update
3. ASM inserts `mfence` before notify
4. Device reads after notify sees consistent data

5.5.2 Device → Rust Visibility

GUARANTEE VIS-2: Device Write Visibility to Rust

GIVEN:

Device writes data to DMA buffer (RX)
Device updates used ring

THEN:

After `asm_vq_poll_rx()` returns, Rust sees device's writes

MECHANISM:

1. Buffer is UC/WC → device writes go to memory
2. ASM reads used.idx (volatile)
3. ASM inserts lfence before reading ring entry
4. ASM inserts lfence before returning buffer index
5. Rust reads buffer after poll returns → sees device data

5.6 Completion Detection

5.6.1 TX Completion Detection

TX completion is detected by comparing used.idx against last_used_idx:

DETECTION TX-COMPLETE:

last_used_idx: Driver's record of last processed used entry
 used.idx: Device's counter of completed descriptors

COMPLETION AVAILABLE iff:

used.idx != last_used_idx
 (accounting for 16-bit wraparound)

ENTRIES AVAILABLE:

count = (used.idx - last_used_idx) & 0xFFFF

5.6.2 RX Completion Detection

Identical mechanism to TX:

DETECTION RX-COMPLETE:

COMPLETION AVAILABLE iff:
 used.idx != last_used_idx

ON EACH COMPLETION:

Buffer contains: 12-byte VirtIO header + Ethernet frame
 Length field: Total bytes written by device
 Payload: buffer[12..length]

5.7 Ordering Violations and Consequences

Violation	Consequence	Detection
Missing sfence before avail.idx	Device sees stale descriptors	Corrupted packets, hangs
Missing mfence before notify	Device notified before data visible	Packet loss, corruption

Violation	Consequence	Detection
Missing lfence after used.idx	Rust reads stale ring entry	Wrong buffer returned
Reading DEVICE-OWNED buffer	Data race with DMA	Corruption, unpredictable
WB memory without clflush	Device reads stale cache data	Corrupted TX packets

5.8 Barrier Summary Table

Operation	Before	After	Rationale
Write descriptor	—	sfence	Visible before idx
Write avail.ring	—	sfence	Visible before idx
Write avail.idx	sfence	mfence	Visible before notify
Write notify	mfence	—	All data visible
Read used.idx	—	lfence	Before ring read
Read used.ring	lfence	lfence	Before buffer read
Read buffer	lfence	—	After ring read

5.9 ASM Implementation Requirements

5.9.1 Submit Implementation Pattern

```
; asm_vq_submit_tx implementation pattern
asm_vq_submit_tx:
    ; ... setup, get descriptor ...

    ; Write descriptor fields
    mov [desc + VIRTQ_DESC_ADDR], buffer_addr
    mov [desc + VIRTQ_DESC_LEN], length
    mov [desc + VIRTQ_DESC_FLAGS], flags

    ; BARRIER: Ensure descriptor visible
    sfence

    ; Write to available ring
    mov [avail_ring + offset], desc_index

    ; BARRIER: Ensure ring entry visible
    sfence

    ; Increment available index
    inc word [avail + VIRTQ_AVAIL_IDX]

    ; BARRIER: Full fence before notify
    mfence
```

```
; Notify device (if needed)
; ... check suppression, write notify register ...

ret
```

5.9.2 Poll Implementation Pattern

```
; asm_vq_poll_rx implementation pattern
asm_vq_poll_rx:
    ; Read used index (volatile)
    mov ax, [used + VIRTQ_USED_IDX]

    ; Compare with last seen
    cmp ax, [state + LAST_USED_IDX]
    je .no_completion

    ; BARRIER: Ensure idx read before ring read
    lfence

    ; Read used ring entry
    ; ... calculate offset, read id and len ...

    ; BARRIER: Ensure ring read before buffer access
    lfence

    ; Update last_used_idx
    inc word [state + LAST_USED_IDX]

    ; Return success with buffer info
    mov eax, 1
    ret

.no_completion:
    xor eax, eax
    ret
```

Chunk 6 — Time & Progress Guarantees

6.1 Time Source

6.1.1 Sole Time Source: TSC

The **only** time source available in this ABI is the Time Stamp Counter (TSC).

Property	Value	Source
Instruction	<code>rdtsc / rdtscp</code>	Intel SDM

Property	Value	Source
Resolution	Sub-nanosecond (at multi-GHz)	CPU-dependent
Width	64 bits	Wraps at ~200 years @ 3GHz
Access	Ring 0 or Ring 3	Always available

No other time sources:

- No PIT (8254) — requires I/O ports, UEFI may have reconfigured
- No HPET — requires discovery, mapping
- No APIC timer — requires interrupt setup
- No RTC — not useful for sub-second timing

Source: REDESIGN §7.4, AUDIT §2.5 TIME-1

6.1.2 TSC Requirements

Before using TSC, the following MUST be verified:

REQUIREMENT TSC-1: Invariant TSC Verification

CHECK: CPUID.80000007H:EDX[bit 8] == 1

IF NOT SET:

TSC frequency may vary with power states
System is non-compliant with this ABI
Behavior: Halt with error

IF SET:

TSC increments at constant rate
Independent of P-states, C-states
Safe for timing measurements

Source: AUDIT §7.1.4

6.1.3 TSC Calibration

TSC frequency MUST be calibrated at boot. Hardcoded frequencies are forbidden.

REQUIREMENT TSC-2: Calibration at Boot

TIMING: Before ExitBootServices (UEFI services available)

METHOD:

1. Read TSC: start = rdtsc()
2. Call UEFI Stall(1_000_000) – 1 second delay
3. Read TSC: end = rdtsc()
4. Frequency = end - start (ticks per second)

STORAGE:

Store in BootHandoff.tsc_freq
Pass to all timing-dependent code

ACCURACY:

$\pm 5\%$ acceptable for network protocol timeouts
UEFI Stall() accuracy varies by firmware

Source: REDESIGN §5.4, §7.6, AUDIT §2.5 TIME-2

6.2 Monotonicity Guarantee

6.2.1 Single-Core Monotonicity

GUARANTEE MONO-1: TSC Monotonicity (Single Core)**GIVEN:**

Invariant TSC verified (CPUID check)
Single-core execution (no SMP)

THEN:

For any two reads t_1, t_2 where t_1 happens-before t_2 :
 $t_2 \geq t_1$ (no backward jumps)

CAVEAT:

Multi-core TSCs may be unsynchronized
This ABI runs single-core only, so this is not a concern

Source: AUDIT §7.1.5

6.2.2 Wraparound Handling

GUARANTEE MONO-2: Wraparound Safety

TSC WIDTH: 64 bits
WRAP PERIOD: ~200 years at 3GHz

HANDLING:

Use wrapping arithmetic for differences:
`elapsed = now.wrapping_sub(start)`

This correctly handles wrap (which won't occur in practice)

6.3 Resolution Guarantee

6.3.1 TSC Resolution

GUARANTEE RES-1: TSC Resolution

MINIMUM RESOLUTION: 1 TSC tick

TYPICAL RESOLUTION: ~0.3-0.5 nanoseconds (at 2-3 GHz)

FOR NETWORK PROTOCOLS:

smoltcp requires millisecond timestamps

1 ms = ~3,000,000 ticks at 3GHz

Resolution is more than sufficient

6.3.2 Conversion Functions

CONVERSION TIME-1: TSC to Milliseconds

FORMULA:

```
milliseconds = tsc_ticks / (tsc_freq / 1000)
```

OR (avoiding overflow):

```
milliseconds = tsc_ticks * 1000 / tsc_freq
```

IMPLEMENTATION:

```
// Safe for tsc_ticks < 2^54 (centuries at 3GHz)
fn tsc_to_ms(ticks: u64, freq: u64) -> u64 {
    ticks / (freq / 1000)
}
```

CONVERSION TIME-2: Milliseconds to TSC

FORMULA:

```
tsc_ticks = milliseconds * (tsc_freq / 1000)
```

IMPLEMENTATION:

```
fn ms_to_tsc(ms: u64, freq: u64) -> u64 {
    ms * (freq / 1000)
}
```

6.4 smoltcp Integration

6.4.1 Timestamp Requirements

smoltcp's **Instant** type uses milliseconds internally:

```
// smoltcp timestamp creation
let timestamp = smoltcp::time::Instant::from_millis(ms as i64);
```

Contract requirement:

- Rust reads TSC via `asm_tsc_read()`
- Rust converts to milliseconds using calibrated `tsc_freq`
- Rust passes milliseconds to smoltcp

6.4.2 Main Loop Timestamp Pattern

```
// CORRECT: Main loop timestamp usage
fn main_loop(tsc_freq: u64) {
    loop {
        // Read TSC once per iteration
        let now_tsc = unsafe { asm_tsc_read() };

        // Convert to smoltcp timestamp
        let now_ms = now_tsc / (tsc_freq / 1000);
        let timestamp = Instant::from_millis(now_ms as i64);

        // Pass to smoltcp
        iface.poll(timestamp, &mut device, &mut sockets);

        // Use same TSC value for timeout checks
        app_state.step(now_tsc, tsc_freq);
    }
}
```

6.4.3 Timer Precision for TCP

TCP Timer	Typical Value	TSC Precision	Adequate?
RTO (min)	200ms	0.0003ms	Yes
RTO (max)	120s	0.0003ms	Yes
Keepalive	2h	0.0003ms	Yes
TIME-WAIT	2×MSL (120s)	0.0003ms	Yes

All TCP timers have precision far exceeding requirements.

6.5 Progress Guarantees

6.5.1 Non-Blocking Function Guarantee

GUARANTEE PROG-1: Bounded Function Execution

ALL ASM interface functions (except `asm_nic_reset`) MUST:

- Return in bounded time
- Not spin-wait on external conditions
- Not loop waiting for device state

```
BOUNDED TIME DEFINITIONS:
asm_tsc_read: < 100 cycles
asm_tsc_read_serialized: < 300 cycles
asm_vq_submit_*: < 200 cycles
asm_vq_poll_*: < 100 cycles
asm_vq_notify: < 1000 cycles (MMIO latency)
asm_bar_*: < 100 cycles
asm_mmio_*: < 500 cycles
```

Source: REDESIGN §2.2 EXEC-3

6.5.2 Exception: Device Reset

EXCEPTION PROG-2: Device Reset Timeout

```
asm_nic_reset() MAY block for bounded time:
MAXIMUM: 100 milliseconds
MECHANISM: Poll status register with TSC timeout
```

IMPLEMENTATION:

```
Write 0 to status register
Loop:
    Read status register
    If status == 0: success, return 0
    If TSC elapsed > 100ms: timeout, return 1
    Brief pause (rep nop)
```

6.5.3 Main Loop Progress

GUARANTEE PROG-3: Main Loop Iteration Bound

```
EACH main loop iteration MUST complete within:
TARGET: 1 millisecond
MAXIMUM: 5 milliseconds (under load)
```

MECHANISM:

- RX poll budget: 32 packets max
- TX submit budget: 32 packets max
- smoltcp.poll(): bounded by socket count and packet count
- App state: O(1) state machine step

Source: REDESIGN §7.1, AUDIT §2.8 LOOP-1

6.6 Timeout Implementation Pattern

6.6.1 Timeout Configuration

```

/// Timeout configuration derived from calibrated TSC
pub struct TimeoutConfig {
    tsc_freq: u64,
}

impl TimeoutConfig {
    pub fn new(tsc_freq: u64) -> Self {
        Self { tsc_freq }
    }

    /// Convert milliseconds to TSC ticks
    #[inline]
    pub fn ms_to_ticks(&self, ms: u64) -> u64 {
        ms * (self.tsc_freq / 1000)
    }

    // Protocol-specific timeouts
    pub fn dhcp_discover(&self) -> u64 { self.ms_to_ticks(5_000) } // 5s
    pub fn dhcp_request(&self) -> u64 { self.ms_to_ticks(3_000) } // 3s
    pub fn tcp_connect(&self) -> u64 { self.ms_to_ticks(30_000) } // 30s
    pub fn dns_query(&self) -> u64 { self.ms_to_ticks(5_000) } // 5s
    pub fn http_response(&self) -> u64 { self.ms_to_ticks(60_000) } // 60s
}

```

6.6.2 Timeout Check Pattern

```

/// Check if timeout has elapsed (non-blocking)
#[inline]
fn is_timed_out(start_tsc: u64, timeout_ticks: u64) -> bool {
    let now = unsafe { asm_tsc_read() };
    now.wrapping_sub(start_tsc) > timeout_ticks
}

// Usage in state machine
impl DhcpState {
    fn step(&mut self, now_tsc: u64, timeouts: &TimeoutConfig) {
        match self {
            DhcpState::Discovering { start_tsc } => {
                if now_tsc.wrapping_sub(*start_tsc) >
timeouts.dhcp_discover() {
                    *self = DhcpState::Failed(Error::Timeout);
                }
                // ... check for response ...
            }
            // ...
        }
    }
}

```

6.7 Forbidden Patterns

6.7.1 Spin-Wait (FORBIDDEN)

```
// FORBIDDEN: Spin-wait loop
fn wait_for_tx_complete(state: &mut VqState) {
    loop { // ← FORBIDDEN
        if let Some(_) = unsafe { asm_vq_poll_tx_complete(state) } {
            return;
        }
        // Spinning forever if device never completes
    }
}
```

6.7.2 Busy Delay (FORBIDDEN)

```
// FORBIDDEN: TSC-based busy delay
fn tsc_delay_us(us: u64) {
    let start = unsafe { asm_tsc_read() };
    let ticks = us * TSC_FREQ / 1_000_000; // Also forbidden: hardcoded
    freq
    while unsafe { asm_tsc_read() } - start < ticks {
        // Busy waiting - FORBIDDEN
    }
}
```

6.7.3 Hardcoded Frequencies (FORBIDDEN)

```
// FORBIDDEN: Hardcoded TSC frequency
const TSC_FREQ: u64 = 2_500_000_000; // WRONG

// CORRECT: Use calibrated value from BootHandoff
fn get_tsc_freq(handoff: &BootHandoff) -> u64 {
    handoff.tsc_freq // Calibrated at boot
}
```

6.8 Time Guarantees Summary

ID	Guarantee	Source
TIME-G1	TSC is sole time source	Contract §6.1
TIME-G2	Invariant TSC verified at boot	AUDIT §7.1.4
TIME-G3	TSC frequency calibrated, never hardcoded	REDESIGN §7.6
TIME-G4	TSC monotonic on single core	AUDIT §7.1.5

ID	Guarantee	Source
TIME-G5	All ASM functions return in bounded cycles	REDESIGN §2.2
TIME-G6	No spin-waits or busy delays	REDESIGN §1.5
TIME-G7	Main loop iteration < 5ms	REDESIGN §7.1
TIME-G8	Timeout checks are non-blocking	Contract §6.6

Chunk 7 — Error Semantics & Failure Modes

7.1 Error Classification

Errors in this ABI are classified into two categories with fundamentally different handling requirements:

7.1.1 Retryable Errors

Retryable errors indicate temporary conditions that may succeed on retry.

Error	Cause	Caller Action
Queue full	All descriptors in use	Retry after collecting completions
No RX packet	RX queue empty	Normal; poll again next iteration
No TX completion	Device hasn't finished	Normal; poll again next iteration
DHCP timeout	No DHCP response	Retry with exponential backoff
TCP connection refused	Remote rejected	Report to app, may retry
DNS timeout	No DNS response	Try fallback server

Caller obligation: May retry immediately or with backoff; must not loop indefinitely.

7.1.2 Fatal Errors

Fatal errors indicate unrecoverable conditions. The system cannot continue.

Error	Cause	System Response
No NIC found	Hardware not present	Halt
VirtIO feature negotiation failed	Incompatible device	Halt
DMA allocation failed	No memory	Halt
ExitBootServices failed	UEFI error	Halt
Device reset timeout	Hardware unresponsive	Halt
Invariant TSC not available	Incompatible CPU	Halt

Error	Cause	System Response
Invalid BootHandoff	Corruption	Halt

System obligation: Log error (if possible), halt processor.

Source: AUDIT §2.9 ERR-1

7.2 ASM Function Error Returns

7.2.1 Error Return Convention

ASM functions use the following error return convention:

Return Value	Meaning
0	Success
1	Retryable error (queue full, not ready, etc.)
0xFFFFFFFF	Sentinel for "no result" (polling functions)
Other	Reserved for future error codes

7.2.2 Function-Specific Error Returns

asm_vq_submit_tx / asm_vq_submit_rx

Return	Meaning	Caller Action
0	Success, buffer submitted	Buffer now DEVICE-OWNED
1	Queue full	Retry after poll_complete

```
// Usage pattern
loop {
    for _ in 0..TX_BUDGET {
        match tx_queue.pop() {
            Some(pkt) => {
                let result = unsafe { asm_vq_submit_tx(state, pkt.idx,
pkt.len) };
                if result != 0 {
                    // Queue full - put packet back, stop for this
                    iteration
                    tx_queue.push_front(pkt);
                    break;
                }
            }
            None => break,
        }
    }
    // Poll completions before next iteration
}
```

asm_vq_poll_tx_complete

Return	Meaning	Caller Action
0..0xFFFFE	Buffer index of completed TX	Return buffer to pool
0xFFFFFFFF	No completion available	Normal; try again later

asm_vq_poll_rx

Return	Meaning	Caller Action
1	Packet received (result populated)	Process packet
0	No packet available	Normal; try again later

asm_nic_reset

Return	Meaning	Caller Action
0	Reset successful	Continue initialization
1	Reset timeout (100ms)	FATAL: Halt

asm_nic_read_mac

Return	Meaning	Caller Action
0	MAC read successfully	Use MAC from output buffer
1	MAC not available	Generate local MAC or fail

7.3 Caller Obligations

7.3.1 Obligation on Retryable Error

OBLIGATION RETRY-1: Handling Retryable Errors

ON receiving retryable error:

1. DO NOT loop immediately retrying
2. DO continue with other main loop phases
3. DO retry on next main loop iteration
4. DO track retry count for eventual timeout

RATIONALE:

Immediate retry loops violate PROG-1 (bounded execution)
Main loop progress ensures completions are collected

7.3.2 Obligation on Fatal Error

```
OBLIGATION FATAL-1: Handling Fatal Errors

ON receiving fatal error:
1. DO NOT retry
2. DO NOT continue execution
3. DO log error if output available (serial, framebuffer)
4. DO halt processor (cli; hlt)

IMPLEMENTATION:
fn fatal_error(code: u32, msg: &str) -> ! {
    // Attempt to log (may fail, that's okay)
    if let Some(output) = debug_output() {
        output.write(msg);
    }
    // Halt
    loop {
        unsafe { asm!("cli; hlt"); }
    }
}
```

7.3.3 Obligation on Polling "Empty"

```
OBLIGATION POLL-1: Handling Empty Poll

ON receiving "no result" from poll function:
1. This is NORMAL operation, not an error
2. DO continue with other main loop phases
3. DO poll again on next iteration
4. DO NOT log or count as error
```

7.4 Error Propagation

7.4.1 ASM → Rust Propagation

ASM communicates errors to Rust via return values only:

```
RULE ERR-PROP-1: ASM Error Returns

ASM functions:
- Return error codes in RAX
- Do NOT modify errno or global state
- Do NOT panic or halt (except internal assertions)
- Do NOT call Rust error handlers

Rust responsibility:
- Check return value after every ASM call
```

- Convert to Rust Result<T, E> if desired
- Decide on retry vs. fatal based on context

7.4.2 Rust → smoltcp Propagation

RULE ERR-PROP-2: Device Trait Error Handling

smoltcp Device trait methods return Option:
 receive() -> Option<(RxToken, TxToken)>
 transmit() -> Option<TxToken>

MAPPING:

- ASM poll_rx returns 0 → receive() returns None
- ASM poll_rx returns 1 → receive() returns Some(...)
- ASM submit_tx returns 1 (full) → transmit() returns None
- ASM submit_tx returns 0 → transmit() returns Some(...)

smoltcp handles None gracefully:

- Stops TX for this poll() call
- Retries on next poll() call

7.5 Panic / Halt Rules

7.5.1 ASM Panic Behavior

RULE PANIC-ASM-1: ASM Must Not Panic

ASM code MUST NOT:

- Call Rust panic machinery
- Trigger CPU exceptions (except unrecoverable)
- Enter infinite loops without timeout

ASM code MAY:

- Return error codes
- Execute hlt on truly unrecoverable conditions
- Trigger debug breakpoint (int3) in debug builds

7.5.2 Rust Panic Behavior

RULE PANIC-RUST-1: Rust Panic Handling

In this no_std environment:

- Default panic handler halts
- No unwinding (panic = abort)
- No recovery from panic

Panic triggers:

- Array bounds check failure
- Integer overflow (debug builds)
- Explicit panic!() calls
- Unwrap on None/Err

RECOMMENDATION:

- Use checked operations
- Handle all Results explicitly
- Avoid unwrap()/expect() in critical paths

7.5.3 Halt Implementation

RULE HALT-1: Processor Halt Sequence

ON unrecoverable error:

```
; Disable interrupts (already disabled, but defensive)
cli

; Halt loop (handles spurious wake)
.halt_loop:
    hlt
    jmp .halt_loop
```

RATIONALE:

- cli ensures no interrupt wake
- hlt stops processor execution
- Loop handles NMI or SMI wake

7.6 Resource Leak Prevention

7.6.1 Leak-Free Error Paths

OBLIGATION LEAK-1: No Resource Leaks on Error

ON any error path, ensure:

- [] All DRIVER-OWNED buffers either:
 - Freed to pool, OR
 - Submitted to queue (becoming DEVICE-OWNED)
- [] All descriptors either:
 - In free list, OR
 - In available ring (pending)
- [] No dangling pointers stored

MECHANISM:

- Use RAII patterns in Rust (Drop trait)
- Track all allocations
- Error paths must clean up

Source: AUDIT §2.9 ERR-2

7.6.2 Example: TX Error Recovery

```
fn try_transmit(driver: &mut VirtioDriver, packet: &[u8]) -> Result<(), TxError> {
    // Allocate buffer (may fail)
    let buffer_idx = driver.alloc_tx_buffer()
        .ok_or(TxError::NoBuffer)?;

    // At this point, buffer is DRIVER-OWNED and we own it

    // Prepare buffer
    let buffer = driver.get_tx_buffer(buffer_idx);
    buffer[..12].fill(0); // VirtIO header
    buffer[12..12+packet.len()].copy_from_slice(packet);

    // Submit (may fail)
    let result = unsafe {
        asm_vq_submit_tx(driver.tx_state, buffer_idx, (12 + packet.len())
as u16)
    };

    if result != 0 {
        // Queue full - we still own the buffer, must free it
        driver.free_tx_buffer(buffer_idx);
        return Err(TxError::QueueFull);
    }

    // Success - buffer is now DEVICE-OWNED
    Ok(())
}
```

7.7 Timeout as Error

7.7.1 Timeout Detection

RULE TIMEOUT-1: Timeout Is Retryable Error

Protocol timeouts (DHCP, TCP, DNS) are retryable errors:

- State machine transitions to "Failed" or "Retry" state
- May retry with backoff
- Eventually becomes fatal if max retries exceeded

Hardware timeouts (device reset) are fatal errors:

- Device is non-responsive
- Cannot continue
- Halt system

7.7.2 Timeout Backoff

```

/// Exponential backoff for retryable timeouts
struct RetryState {
    attempts: u32,
    max_attempts: u32,
    base_timeout_ms: u64,
    max_timeout_ms: u64,
}

impl RetryState {
    fn next_timeout(&mut self, timeouts: &TimeoutConfig) -> Option<u64> {
        if self.attempts >= self.max_attempts {
            return None; // Give up
        }

        // Exponential backoff: base * 2^attempts, capped at max
        let timeout_ms = self.base_timeout_ms
            .saturating_mul(1 << self.attempts)
            .min(self.max_timeout_ms);

        self.attempts += 1;
        Some(timeouts.ms_to_ticks(timeout_ms))
    }
}

```

7.8 Error Codes Summary

Code	Name	Category	Meaning
0	SUCCESS	—	Operation completed
1	QUEUE_FULL	Retryable	Virtqueue has no free descriptors
1	TIMEOUT	Retryable/Fatal	Operation timed out
1	NOT_AVAILABLE	Retryable	Resource not ready
0xFFFFFFFF	NO_RESULT	—	Poll found nothing (normal)

7.9 Error Invariants Summary

ID	Invariant	Consequence of Violation
ERR-INV-1	Retryable errors must not cause infinite retry loops	System hang
ERR-INV-2	Fatal errors must halt system	Continued corruption
ERR-INV-3	All error paths must free owned resources	Memory/descriptor leak
ERR-INV-4	ASM returns error codes, does not halt	Rust loses control

ID	Invariant	Consequence of Violation
ERR-INV-5	Rust checks all ASM return values	Silent failures

Chunk 8 — Safety, Undefined Behavior & Misuse

8.1 Undefined Behavior Definition

Undefined Behavior (UB) in this contract means the system may exhibit any behavior, including but not limited to:

- Data corruption
- Security vulnerabilities
- System hang
- Incorrect results
- Apparently correct operation (most dangerous)

Philosophy: This contract adopts a "hostile misuse" stance. Any deviation from the contract is UB, even if it "seems to work" in testing.

8.2 Explicit UB Conditions

8.2.1 Calling Convention Violations

Violation	Why It's UB
Wrong parameter registers	Function receives garbage values
Missing shadow space	Stack corruption
Misaligned stack	Potential crash, corruption
Wrong return register	Caller receives garbage
Clobbering non-volatile registers	Caller state destroyed
Non-zero direction flag	String operations corrupted

8.2.2 Ownership Violations

Violation	Why It's UB
Reading DEVICE-OWNED buffer	Data race with DMA engine
Writing DEVICE-OWNED buffer	Data race with DMA engine
Double-free of buffer	Allocator corruption
Use-after-free	Dangling pointer access
Submitting FREE buffer	Invalid descriptor

Violation	Why It's UB
Submitting already-submitted buffer	Descriptor reuse conflict

8.2.3 Memory Ordering Violations

Violation	Why It's UB
Missing sfence before avail.idx	Device sees stale descriptors
Missing mfence before notify	Notify races with data
Missing lfence after used.idx	Reading stale completion
Compiler reordering of DMA access	Device sees wrong data
Using WB memory without clflush	Cache incoherence

8.2.4 Pointer Violations

Violation	Why It's UB
Null pointer to ASM function	Dereferencing null
Misaligned pointer	Alignment fault
Out-of-bounds pointer	Memory corruption
Pointer to freed memory	Use-after-free
Pointer outside DMA region	Invalid DMA address

8.2.5 State Violations

Violation	Why It's UB
Using NIC before initialization	Hardware in unknown state
Using queues before setup	Invalid ring pointers
Calling EBS functions post-EBS	Firmware not available
Enabling interrupts	No handler installed
Running on AP core	TSC unsynchronized, races

8.2.6 Timing Violations

Violation	Why It's UB
Hardcoded TSC frequency	Wrong timeout calculations
Using TSC without invariant check	Frequency may vary
Infinite spin-wait	System hang

Violation	Why It's UB
Main loop > 5ms	Protocol timeouts, poor performance

8.3 Rust **unsafe** Boundaries

8.3.1 Required **unsafe** Blocks

All ASM function calls require **unsafe**:

```
// ALL ASM calls must be in unsafe blocks
let tsc = unsafe { asm_tsc_read() };

let result = unsafe {
    asm_vq_submit_tx(state, idx, len)
};

unsafe {
    asm_bar_mfence();
}
```

Rationale: The compiler cannot verify:

- Pointer validity
- Ownership rules
- Memory ordering
- State preconditions

8.3.2 **unsafe** Contract Documentation

Each **unsafe** block SHOULD document what invariants the caller guarantees:

```
// SAFETY:
// - state points to valid, initialized VirtqueueState
// - buffer_idx < queue_size
// - Buffer at buffer_idx is DRIVER-OWNED
// - Buffer content fully written (including 12-byte header)
let result = unsafe {
    asm_vq_submit_tx(state, buffer_idx, length)
};
```

8.3.3 **unsafe** Trait Implementations

The smoltcp Device trait implementation requires unsafe operations:

```
impl smoltcp::phy::Device for VirtioDevice {
    type RxToken<'a> = VirtioRxToken<'a>;
```

```

type TxToken<'a> = VirtioTxToken<'a>;

fn receive(&mut self, _timestamp: Instant) ->
Option<(Self::RxToken<'_>, Self::TxToken<'_>) {
    let mut result = RxResult::default();

    // SAFETY:
    // - self.rx_state is valid VirtqueueState
    // - result is valid output buffer
    let status = unsafe {
        asm_vq_poll_rx(&mut self.rx_state, &mut result)
    };

    if status == 0 {
        return None;
    }

    // Buffer at result.buffer_index is now DRIVER-OWNED
    Some((
        VirtioRxToken {
            device: self,
            buffer_idx: result.buffer_index,
            length: result.length,
        },
        VirtioTxToken { device: self },
    ))
}

// ... transmit() similarly
}

```

8.4 Contract Violation Consequences

8.4.1 Severity Levels

Level	Example	Consequence
Catastrophic	DMA to wrong address	System corruption, potential security hole
Severe	Ownership violation	Data corruption, packet loss
Moderate	Missing barrier	Intermittent corruption
Minor	Wrong error handling	Suboptimal behavior

8.4.2 Detection Difficulty

Violation Type	Detection	Notes
Calling convention	Immediate crash or corruption	Usually obvious
Ownership	Intermittent corruption	Hard to reproduce

Violation Type	Detection	Notes
Memory ordering	Rare corruption	May only appear under load
Pointer validity	Crash or silent corruption	May work in testing
Timing	Timeouts, performance issues	May seem like network issues

8.4.3 QEMU vs. Real Hardware

Warning: Many violations that "work" in QEMU will fail on real hardware:

Violation	QEMU Behavior	Real Hardware
Missing barriers	Usually works	Corruption
WB memory for DMA	Works (coherent emulation)	Corruption
Timing assumptions	May work	Fails under load
Ownership races	Hidden by synchronous emulation	Data corruption

8.5 Defensive Programming Requirements

8.5.1 Pointer Validation (Debug Builds)

```
#[cfg(debug_assertions)]
fn validate_vq_state(state: *const VirtqueueState) {
    assert!(!state.is_null(), "VirtqueueState pointer is null");
    assert!(state as usize % 64 == 0, "VirtqueueState not cache-aligned");

    let state_ref = unsafe { &*state };
    assert!(state_ref.queue_size.is_power_of_two(), "Queue size not power
of 2");
    assert!(state_ref.queue_size <= 32768, "Queue size exceeds max");
}
```

8.5.2 Bounds Checking

```
fn get_buffer(&self, idx: u16) -> &mut [u8] {
    assert!((idx as usize) < self.buffer_count, "Buffer index out of
bounds");
    // ...
}

fn submit_tx(&mut self, idx: u16, len: u16) -> Result<(), Error> {
    if idx as usize >= self.queue_size as usize {
        return Err(Error::InvalidIndex);
    }
    if len as usize > self.buffer_size {
        return Err(Error::BufferTooLarge);
    }
}
```

```

    }
    // ...
}

```

8.5.3 State Assertions

```

impl VirtioDriver {
    fn assert_initialized(&self) {
        debug_assert!(self.initialized, "Driver used before
initialization");
        debug_assert!(self.status & STATUS_DRIVER_OK != 0, "Device not
ready");
    }

    pub fn submit_tx(&mut self, ...) -> Result<...> {
        self.assert_initialized();
        // ...
    }
}

```

8.6 Misuse Scenarios

8.6.1 Scenario: Accessing DEVICE-OWNED Buffer

```

// MISUSE: Reading buffer after submit
let idx = driver.alloc_tx_buffer();
let buffer = driver.get_buffer(idx);
buffer.copy_from_slice(&packet);

unsafe { asm_vq_submit_tx(state, idx, len); }

// BUG: Buffer is now DEVICE-OWNED
let first_byte = buffer[0]; // UB! Data race with DMA

```

Consequence: If device is simultaneously transmitting, read may see partial/corrupted data.

8.6.2 Scenario: Missing Memory Barrier

```

// MISUSE: Manual descriptor update without barriers
unsafe {
    (*desc).addr = buffer_addr;
    (*desc).len = packet_len;
    (*desc).flags = 0;

    // BUG: Missing sfence
}

```

```

(*avail).ring[idx] = desc_idx;
(*avail).idx += 1;

// BUG: Missing mfence

asm_mmio_write16(notify_addr, queue_idx);
}

```

Consequence: Device may see stale descriptor, transmit garbage or crash.

8.6.3 Scenario: Spin-Wait in Callback

```

impl TxToken for VirtioTxToken<'_> {
    fn consume<R, F>(&self, len: usize, f: F) -> R
    where F: FnOnce(&&mut [u8]) -> R
    {
        let idx = self.driver.alloc_tx_buffer();
        let buffer = self.driver.get_buffer(idx);
        buffer[..12].fill(0);
        let result = f(&&mut buffer[12..12+len]);

        unsafe { asm_vq_submit_tx(self.driver.tx_state, idx, (12+len) as
u16) };

        // BUG: Waiting for completion inside callback
        loop {
            if unsafe { asm_vq_poll_tx_complete(self.driver.tx_state) } != 0xFFFFFFFF {
                break;
            }
        } // Infinite loop if device slow/stuck

        result
    }
}

```

Consequence: System hangs, violates progress guarantee.

8.6.4 Scenario: Hardcoded TSC Frequency

```

// MISUSE: Hardcoded frequency
const TSC_FREQ: u64 = 2_500_000_000; // "It's always 2.5GHz, right?"

fn timeout_5_seconds() -> u64 {
    5 * TSC_FREQ // WRONG on any other CPU
}

```

Consequence: On 3.5GHz CPU, timeout is 3.5 seconds. On 2.0GHz CPU, timeout is 6.25 seconds.

8.7 Formal Precondition/Postcondition Summary

8.7.1 `asm_vq_submit_tx`

PRECONDITIONS:

- P1: `vq_state` != `NULL`
- P2: `vq_state` points to valid, initialized `VirtqueueState`
- P3: `buffer_index` < `vq_state.queue_size`
- P4: `Buffer[buffer_index]` is DRIVER-OWNED
- P5: Buffer content is fully written
- P6: `buffer_length` ≤ `vq_state.buffer_size`
- P7: `buffer_length` ≥ 12 (minimum VirtIO header)

POSTCONDITIONS (if return == 0):

- Q1: `Buffer[buffer_index]` is DEVICE-OWNED
- Q2: Descriptor submitted to available ring
- Q3: Device notified (if appropriate)

POSTCONDITIONS (if return == 1):

- Q4: `Buffer[buffer_index]` is still DRIVER-OWNED
- Q5: No state changed

8.7.2 `asm_vq_poll_rx`

PRECONDITIONS:

- P1: `vq_state` != `NULL`
- P2: `vq_state` points to valid, initialized `VirtqueueState`
- P3: `result` != `NULL`
- P4: `result` points to valid `RxResult`

POSTCONDITIONS (if return == 1):

- Q1: `result.buffer_index` contains valid index
- Q2: `Buffer[result.buffer_index]` is DRIVER-OWNED
- Q3: Buffer contains packet data (12B header + frame)
- Q4: `result.length` contains actual bytes

POSTCONDITIONS (if return == 0):

- Q5: `result` contents are undefined
- Q6: No buffers changed ownership

8.8 Safety Invariants Summary

ID	Invariant	Category
SAFE-1	All ASM calls in <code>unsafe</code> blocks	Rust safety
SAFE-2	Pointer validity documented in <code>// SAFETY:</code>	Documentation
SAFE-3	No access to DEVICE-OWNED buffers	Ownership

ID	Invariant	Category
SAFE-4	All barriers placed per contract	Memory ordering
SAFE-5	No spin-waits or unbounded loops	Progress
SAFE-6	TSC frequency from calibration only	Timing
SAFE-7	All return values checked	Error handling
SAFE-8	Debug assertions for invariants	Defensive

Chunk 9 — Final Contract Review & Freeze Statement

9.1 Internal Consistency Verification

This section verifies that all chunks of this contract are internally consistent and do not contradict each other.

9.1.1 Cross-Chunk Consistency Matrix

Chunk	References	Consistent With	Verified
1. Scope	REDESIGN, AUDIT	Chunks 2-8	✓
2. ABI	Intel SDM, MS ABI	Chunk 3 (interface)	✓
3. Interface	Chunk 2 (ABI), Chunk 4 (ownership)	Chunks 5, 6	✓
4. Ownership	REDESIGN §3.3, AUDIT §2.1	Chunk 3, 5	✓
5. Ordering	AUDIT §5.4, Intel SDM	Chunk 3, 4	✓
6. Time	REDESIGN §7, AUDIT §2.5	Chunk 1, 3	✓
7. Errors	AUDIT §2.9	Chunk 3, 4, 8	✓
8. Safety	All chunks	Chunks 1-7	✓

9.1.2 Invariant Consistency

All invariants defined across chunks have been checked for:

- No contradictions between invariants
- No impossible-to-satisfy combinations
- Complete coverage of the ABI surface

Invariant Count Summary:

Category	Count
Execution (EXEC-INV-*)	5
Memory (MEM-INV-*)	6

Category	Count
Timing (TIME-INV-*)	4
VirtIO (VIO-INV-*)	6
Barrier (BAR-INV-*)	4
Lifetime (LIFE-*)	6
Time Guarantee (TIME-G*)	8
Error (ERR-INV-*)	5
Safety (SAFE-*)	8
Total	52

9.1.3 Source Document Traceability

Every normative statement in this contract traces to:

- NETWORK_STACK_REDESIGN.md (v1.0 Frozen), OR
- NETWORK_STACK_AUDIT.md (v1.1), OR
- Intel SDM / VirtIO Spec / UEFI Spec (for hardware/protocol details)

No statement introduces new architectural requirements not present in source documents.

9.2 Completeness Checklist

9.2.1 ABI Coverage

Aspect	Documented	Section
Calling convention	✓	Chunk 2
Parameter passing	✓	§2.1.1
Return values	✓	§2.1.2
Register preservation	✓	§2.2
Stack requirements	✓	§2.1.3, §2.1.4
Clobber rules	✓	§2.4

9.2.2 Interface Coverage

Function	Documented	Section
asm_tsc_read	✓	§3.3.1
asm_tsc_read_serialized	✓	§3.3.1
asm_vq_submit_tx	✓	§3.3.2

Function	Documented	Section
asm_vq_poll_tx_complete	✓	§3.3.2
asm_vq_submit_rx	✓	§3.3.2
asm_vq_poll_rx	✓	§3.3.2
asm_vq_notify	✓	§3.3.2
asm_bar_sfence	✓	§3.3.3
asm_bar_lfence	✓	§3.3.3
asm_bar_mfence	✓	§3.3.3
asm_mmio_read32	✓	§3.3.4
asm_mmio_write32	✓	§3.3.4
asm_mmio_read16	✓	§3.3.4
asm_mmio_write16	✓	§3.3.4
asm_nic_reset	✓	§3.3.5
asm_nic_set_status	✓	§3.3.5
asm_nic_get_status	✓	§3.3.5
asm_nic_read_features	✓	§3.3.5
asm_nic_write_features	✓	§3.3.5
asm_nic_read_mac	✓	§3.3.5

9.2.3 Semantic Coverage

Aspect	Documented	Section
Ownership states	✓	Chunk 4
Ownership transitions	✓	§4.3
Buffer lifecycles	✓	§4.4, §4.5
Memory barriers	✓	Chunk 5
Barrier placement	✓	§5.3
Visibility guarantees	✓	§5.5
Time source	✓	Chunk 6
Progress guarantees	✓	§6.5
Error classification	✓	Chunk 7
Error returns	✓	§7.2

Aspect	Documented	Section
UB conditions	✓	Chunk 8
Safety boundaries	✓	§8.3

9.3 Known Limitations

This contract acknowledges the following limitations:

9.3.1 Scope Limitations

Limitation	Rationale
VirtIO only	Other NICs require separate contracts
Single-core only	Multi-core requires synchronization extensions
No interrupts	Interrupt support requires IDT, ISR contracts
IPv4 only	IPv6 requires additional protocol work
No TLS	Cryptography out of scope for v1

9.3.2 Implementation Guidance Not Provided

This contract does NOT specify:

- Internal ASM register allocation
- Internal Rust data structure layouts (beyond shared structs)
- Optimization strategies
- Debug/logging mechanisms
- Test methodologies

These are implementation concerns, not ABI concerns.

9.4 Versioning Policy

9.4.1 Version Number Scheme

Version: MAJOR.MINOR

MAJOR: Incremented when:

- Any breaking change to function signatures
- Any breaking change to struct layouts
- Any change to calling convention
- Any change that requires both ASM and Rust changes

MINOR: Incremented when:

- New functions added (backward compatible)
- Documentation clarifications
- Bug fixes that don't change interface

9.4.2 Compatibility Rules

RULE COMPAT-1: Major Version Compatibility

Code compiled against contract version X.* is INCOMPATIBLE with code compiled against version Y.* where X ≠ Y.

RULE COMPAT-2: Minor Version Compatibility

Code compiled against contract version X.A is FORWARD COMPATIBLE with code compiled against version X.B where B > A.

(New features may not be available, but existing features work.)

RULE COMPAT-3: Linking Requirement

ASM object files and Rust object files MUST be compiled against the same major version of this contract.

9.4.3 Deprecation Policy

POLICY DEP-1: Function Deprecation

Functions may be deprecated but not removed within a major version.

Deprecation is indicated by:

- Documentation note
- Compiler warning (if tooling supports)

Deprecated functions are removed in the next major version.

9.5 Contract Freeze Declaration

9.5.1 Freeze Statement

CONTRACT FREEZE DECLARATION

Document: NETWORK_ASM_RUST_ABI_CONTRACT.md

Version: 1.0

Date: 2026-01-09

Status: FROZEN

This document represents the frozen v1.0 ABI contract between ASM and Rust for the MorpheusX post-ExitBootServices networking runtime.

All implementations MUST conform to this contract.

Any deviation constitutes undefined behavior.

Modifications to this contract require:

1. Formal review against source documents
2. Impact analysis on existing implementations
3. Version increment per §9.4
4. Explicit sign-off

9.5.2 Authority Chain

This contract is authoritative for:

→ ASM ↔ Rust interface behavior

This contract derives authority from:

→ NETWORK_STACK_REDESIGN.md (v1.0 Frozen)
→ NETWORK_STACK_AUDIT.md (v1.1)

These documents derive authority from:

→ VirtIO Specification 1.2 (OASIS)
→ UEFI Specification 2.10
→ Intel 64 and IA-32 Architectures SDM
→ Microsoft x64 ABI documentation

9.5.3 Conformance Statement

An implementation conforms to this contract if and only if:

1. All ASM functions match the signatures in Chunk 3
2. All ASM functions implement the specified behavior
3. All ASM functions place barriers as specified in Chunk 5
4. All Rust code calls ASM functions with valid preconditions
5. All Rust code handles return values per Chunk 7
6. All ownership rules from Chunk 4 are respected
7. All timing constraints from Chunk 6 are met
8. No undefined behavior conditions from Chunk 8 are triggered

9.6 Future Work (Out of Scope)

The following are explicitly deferred to future contract versions:

Item	Target Version	Notes
Multi-core support	v2.0	Requires synchronization primitives
Interrupt support	v2.0	Requires IDT, ISR interface
Additional NIC drivers	v1.x	Additive, minor version
Hardware offload support	v2.0	Changes buffer semantics
IPv6 support	v1.x	smoltcp supports, additive

9.7 Document End

=====

=====

END OF CONTRACT DOCUMENT

Document: MorpheusX ASM ↔ Rust ABI Contract
 Version: 1.0 (FROZEN)
 Date: 2026-01-09
 Chunks: 9 of 9 complete

Source Documents:

- NETWORK_STACK_REDESIGN.md v1.0 (Frozen)
- NETWORK_STACK_AUDIT.md v1.1

Next Steps (NOT part of this document):

- Binding generation (extern "win64")
- Rust unsafe annotation pass
- ASM enforcement scaffolding